

## SCRIPT PREDICCION PARTE 2

June 8, 2023

```
[1]: import pandas as pd
import numpy as np
import math
import warnings
warnings.filterwarnings("ignore")
```

```
[2]: df = pd.read_pickle('df.pkl')
df
```

```
[2]:
```

	CAMPAÑA	ID_FINCA	ID_ZONA	ID_ESTACION	ALTITUD	VARIEDAD	MOD0	\
0	14	76953	515	4	660.000000	26	2	
1	14	84318	515	4	660.000000	26	2	
2	14	85579	340	4	520.000000	32	2	
3	14	69671	340	4	520.000000	32	2	
4	14	14001	852	14	659.097938	81	1	
...	...	...	...	...	...	...	...	
9596	22	37461	239	6	700.000000	52	2	
9597	22	58769	239	6	700.000000	32	2	
9598	22	58769	239	6	700.000000	59	2	
9599	22	88928	239	6	700.000000	40	2	
9600	22	88928	239	6	700.000000	52	2	

	TIPO	COLOR	SUPERFICIE	PRODUCCION
0	0	1	4.902003	22215.0
1	0	1	4.633973	22215.0
2	0	1	4.995608	20978.0
3	0	1	7.126296	40722.0
4	0	1	4.354864	14126.0
...	...	...	...	...
9596	0	1	3.680000	NaN
9597	0	1	4.250000	NaN
9598	0	1	4.070000	NaN
9599	0	1	4.572700	NaN
9600	0	1	1.609900	NaN

```
[9601 rows x 11 columns]
```

La estrategia que vamos a desarrollar va a ser codificar la variable ID\_ESTACION por la media para

cada estacion de la variable meteorológica que más correlación tenga con la variable PRODUCCION

```
[3]: #Definimos una funcion para meter las variables meteorológicas de interés
def metervariables(df):
    meses=['enero', 'febrero', 'marzo', 'abril', 'mayo', 'junio']
    for m in range(len(meses)):
        precip_col = f"precip_{meses[m]}"
        df[precip_col] = np.nan

    # Rellenamos los valores de precip_mes en función de CAMPAÑA y ID_ESTACION
    for c in range(16, 23):
        for e in range(20):
            for m in range(6):
                precip_mes = globals()[f"precip_mes_{c}"][e][m]
                mes_index = meses[m]
                precip_col = f"precip_{mes_index}"
                df.loc[(df['CAMPAÑA'] == c) & (df['ID_ESTACION'] == e),
↪precip_col] = precip_mes

    for m in range(len(meses)):
        windspeed_col = f"windspeed_{meses[m]}"
        df[windspeed_col] = np.nan

    # Rellenamos los valores de windspeed_mes en función de CAMPAÑA y
↪ID_ESTACION
    for c in range(16, 23):
        for e in range(20):
            for m in range(6):
                windspeed_mes = globals()[f"windspeed_mes_{c}"][e][m]
                mes_index = meses[m]
                windspeed_col = f"windspeed_{mes_index}"
                df.loc[(df['CAMPAÑA'] == c) & (df['ID_ESTACION'] == e),
↪windspeed_col] = windspeed_mes

    # Rellenamos los valores de temp_mes en función de CAMPAÑA y ID_ESTACION
    for m in range(len(meses)):
        temp_col = f"t_mes_{meses[m]}"
        df[temp_col] = np.nan
    for c in range(16, 23):
        for e in range(20):
            for m in range(6):
                temp_mes = globals()[f"t_mes_{c}"][e][m]
                mes_index = meses[m]
                temp_col = f"t_mes_{mes_index}"
                df.loc[(df['CAMPAÑA'] == c) & (df['ID_ESTACION'] == e),
↪temp_col] = temp_mes
    for m in range(len(meses)):
```

```

        rhum_col = f"rhumb_mes_{meses[m]}"
        df[rhum_col] = np.nan
    for c in range(16, 23):
        for e in range(20):
            for m in range(6):
                rhumb_mes = globals()[f"rhumb_mes_{c}"][e][m]
                mes_index = meses[m]
                rhumb_col = f"rhumb_mes_{mes_index}"
                df.loc[(df['CAMPAÑA'] == c) & (df['ID_ESTACION'] == e),
↪rhumb_col] = rhumb_mes
    return df

```

```

[4]: #Importamos las variables
import pickle

with open('rhumb_mes_16.pkl', 'rb') as file:
    rhumb_mes_16 = pickle.load(file)

with open('rhumb_mes_17.pkl', 'rb') as file:
    rhumb_mes_17 = pickle.load(file)

with open('rhumb_mes_18.pkl', 'rb') as file:
    rhumb_mes_18 = pickle.load(file)

with open('rhumb_mes_19.pkl', 'rb') as file:
    rhumb_mes_19 = pickle.load(file)

with open('rhumb_mes_20.pkl', 'rb') as file:
    rhumb_mes_20 = pickle.load(file)

with open('rhumb_mes_21.pkl', 'rb') as file:
    rhumb_mes_21 = pickle.load(file)

with open('rhumb_mes_22.pkl', 'rb') as file:
    rhumb_mes_22 = pickle.load(file)

with open('windspeed_mes_16.pkl', 'rb') as file:
    windspeed_mes_16 = pickle.load(file)

with open('windspeed_mes_17.pkl', 'rb') as file:
    windspeed_mes_17 = pickle.load(file)

with open('windspeed_mes_18.pkl', 'rb') as file:
    windspeed_mes_18 = pickle.load(file)

with open('windspeed_mes_19.pkl', 'rb') as file:
    windspeed_mes_19 = pickle.load(file)

```

```

with open('windspeed_mes_20.pkl', 'rb') as file:
    windspeed_mes_20 = pickle.load(file)

with open('windspeed_mes_21.pkl', 'rb') as file:
    windspeed_mes_21 = pickle.load(file)

with open('windspeed_mes_22.pkl', 'rb') as file:
    windspeed_mes_22 = pickle.load(file)

with open('t_mes_15.pkl', 'rb') as file:
    t_mes_15 = pickle.load(file)

with open('t_mes_16.pkl', 'rb') as file:
    t_mes_16 = pickle.load(file)

with open('t_mes_17.pkl', 'rb') as file:
    t_mes_17 = pickle.load(file)

with open('t_mes_18.pkl', 'rb') as file:
    t_mes_18 = pickle.load(file)

with open('t_mes_19.pkl', 'rb') as file:
    t_mes_19 = pickle.load(file)

with open('t_mes_20.pkl', 'rb') as file:
    t_mes_20 = pickle.load(file)

with open('t_mes_21.pkl', 'rb') as file:
    t_mes_21 = pickle.load(file)

with open('t_mes_22.pkl', 'rb') as file:
    t_mes_22 = pickle.load(file)

with open('precip_mes_16.pkl', 'rb') as file:
    precip_mes_16 = pickle.load(file)

with open('precip_mes_17.pkl', 'rb') as file:
    precip_mes_17 = pickle.load(file)

with open('precip_mes_18.pkl', 'rb') as file:
    precip_mes_18 = pickle.load(file)

with open('precip_mes_19.pkl', 'rb') as file:
    precip_mes_19 = pickle.load(file)

with open('precip_mes_20.pkl', 'rb') as file:

```

```

precip_mes_20 = pickle.load(file)

with open('precip_mes_21.pkl', 'rb') as file:
    precip_mes_21 = pickle.load(file)

with open('precip_mes_22.pkl', 'rb') as file:
    precip_mes_22 = pickle.load(file)

```

```

[5]: #Creamos un dataframe df2 donde metemos las variables meteorológicas
df2=metervariables(df)
df = pd.read_pickle('df.pkl')

```

```

[6]: #Estación 4
#Para la estación 4, no hay informacion para muchas de las variables_
↳ meteorológicas
df2[df2['ID_ESTACION']==4]

```

```

[6]:  CAMPAÑA  ID_FINCA  ID_ZONA  ID_ESTACION  ALTITUD  VARIEDAD  MODO  TIPO  \
0      14      76953      515           4      660.0      26      2      0
1      14      84318      515           4      660.0      26      2      0
2      14      85579      340           4      520.0      32      2      0
3      14      69671      340           4      520.0      32      2      0

    COLOR  SUPERFICIE  ...  t_mes_marzo  t_mes_abril  t_mes_mayo  t_mes_junio  \
0      1      4.902003  ...           NaN           NaN           NaN           NaN
1      1      4.633973  ...           NaN           NaN           NaN           NaN
2      1      4.995608  ...           NaN           NaN           NaN           NaN
3      1      7.126296  ...           NaN           NaN           NaN           NaN

    rhum_mes_enero  rhum_mes_febrero  rhum_mes_marzo  rhum_mes_abril  \
0              NaN              NaN              NaN              NaN
1              NaN              NaN              NaN              NaN
2              NaN              NaN              NaN              NaN
3              NaN              NaN              NaN              NaN

    rhum_mes_mayo  rhum_mes_junio
0              NaN              NaN
1              NaN              NaN
2              NaN              NaN
3              NaN              NaN

```

[4 rows x 35 columns]

```

[7]: #Como son solo 4 observaciones decidimos eliminarlas
df = df.drop(df[df['ID_ESTACION'] == 4].index)
df2 = df2.drop(df2[df2['ID_ESTACION'] == 4].index)

```

```
[8]: #Definimos una función que calcule la media de las variables meteorológicas
      ↪para cada ID_ESTACION en los meses enero-junio
def calcular_variable_media(df,columnas_meses):
    # Filtrar el DataFrame para mantener solo las columnas de los meses de
    ↪enero a junio y la columna 'ID_ESTACION'
    df_filt = df[columnas_meses + ['ID_ESTACION']].copy()

    # Calcular la temperatura media para cada valor de ID_ESTACION
    df_media_temperaturas = df_filt.groupby('ID_ESTACION')[columnas_meses].
    ↪mean()
    df_media_temperaturas = df_media_temperaturas.reset_index()

    return df_media_temperaturas
```

Aplicamos la funcion anterior y vamos metiendo los resultados obtenidos en nuestro dataframe:

```
[9]: df_media=calcular_variable_media(df2, columnas_meses = ['windspeed_enero',
    ↪'windspeed_febrero', 'windspeed_marzo', 'windspeed_abril', 'windspeed_mayo',
    ↪'windspeed_junio'])
```

```
[10]: encoded_id_estacion=df_media.groupby('ID_ESTACION').mean().mean(axis=1).
    ↪reset_index(name='windspeed_media')
encoded_id_estacion
```

```
[10]:
```

	ID_ESTACION	windspeed_media
0	0	8.055134
1	1	13.853971
2	2	10.309313
3	3	12.022968
4	5	10.712688
5	6	11.664805
6	7	11.258397
7	8	12.295432
8	9	10.534987
9	10	10.132563
10	11	10.211159
11	12	11.647173
12	13	12.026949
13	14	12.172118
14	15	11.849458
15	16	10.843103
16	17	12.523165
17	18	10.417981
18	19	10.871163

```
[11]: df['windspeed_media'] = df['ID_ESTACION'].map(encoded_id_estacion.
    ↪set_index('ID_ESTACION')['windspeed_media'])
```

```
[12]: df_media_temperaturas=calcular_variable_media(df2, columnas_meses =
↳ ['t_mes_enero', 't_mes_febrero', 't_mes_marzo', 't_mes_abril', 't_mes_mayo',
↳ 't_mes_junio'])
```

```
[13]: encoded_id_estacion=df_media_temperaturas.groupby('ID_ESTACION').mean().
↳ mean(axis=1).reset_index(name='tmedia')
encoded_id_estacion
```

```
[13]:
```

	ID_ESTACION	tmedia
0	0	14.287836
1	1	11.802972
2	2	13.947971
3	3	12.972376
4	5	12.963303
5	6	12.589099
6	7	13.316537
7	8	12.737542
8	9	13.657491
9	10	14.182774
10	11	13.752891
11	12	13.713871
12	13	13.119273
13	14	12.094942
14	15	12.587852
15	16	12.946614
16	17	12.394114
17	18	12.946488
18	19	13.676968

```
[14]: df['t_media'] = df['ID_ESTACION'].map(encoded_id_estacion.
↳ set_index('ID_ESTACION')['tmedia'])
```

```
[15]: df_media=calcular_variable_media(df2, columnas_meses = ['rhum_mes_enero',
↳ 'rhum_mes_febrero', 'rhum_mes_marzo', 'rhum_mes_abril', 'rhum_mes_mayo',
↳ 'rhum_mes_junio'])
```

```
[16]: encoded_id_estacion=df_media.groupby('ID_ESTACION').mean().mean(axis=1).
↳ reset_index(name='rhum_media')
encoded_id_estacion
```

```
[16]:
```

	ID_ESTACION	rhum_media
0	0	68.650047
1	1	65.857631
2	2	64.302105
3	3	66.280790
4	5	66.633674
5	6	66.438743

6	7	65.982460
7	8	66.031945
8	9	65.956254
9	10	65.434847
10	11	64.647435
11	12	65.781480
12	13	65.082082
13	14	66.492972
14	15	66.012340
15	16	66.640886
16	17	67.463580
17	18	66.513244
18	19	66.100459

```
[17]: df['rhum_media'] = df['ID_ESTACION'].map(encoded_id_estacion.
↳set_index('ID_ESTACION')['rhum_media'])
```

```
[18]: df_media=calcular_variable_media(df2, columnas_meses = ['precip_enero',
↳'precip_febrero', 'precip_marzo', 'precip_abril', 'precip_mayo',
↳'precip_junio'])
```

```
[19]: encoded_id_estacion=df_media.groupby('ID_ESTACION').mean().mean(axis=1).
↳reset_index(name='precip_media')
encoded_id_estacion
```

```
[19]:
```

	ID_ESTACION	precip_media
0	0	65.166667
1	1	55.575000
2	2	39.023481
3	3	40.346064
4	5	41.422958
5	6	55.656785
6	7	39.262180
7	8	38.912033
8	9	40.980680
9	10	49.085008
10	11	40.451190
11	12	40.099020
12	13	39.370690
13	14	40.389625
14	15	38.996225
15	16	41.668350
16	17	35.540517
17	18	44.177987
18	19	39.823795



```
[20]: df['precip_media'] = df['ID_ESTACION'].map(encoded_id_estacion.
      ↪set_index('ID_ESTACION')['precip_media'])
```

Una variable que no incluimos en nuestros modelos anteriores y que podría ser relevante es el índice de radiación ultravioleta. Acudimos a la base de datos de Meteo para obtener información sobre esta variable.

```
[21]: meteo = pd.read_csv('DATOS_METEO.txt', delimiter='|')
```

```
[22]: # Creamos columnas separadas para el año, mes, día y hora
meteo['validTimeUtc'] = pd.to_datetime(meteo['validTimeUtc'])

meteo['año'] = meteo['validTimeUtc'].dt.year
meteo['mes'] = meteo['validTimeUtc'].dt.month
meteo['dia'] = meteo['validTimeUtc'].dt.day
meteo['hora'] = meteo['validTimeUtc'].dt.hour

meteo
```

```
[22]:
```

	validTimeUtc	precip1Hour	precip6Hour	precip24Hour	\
0	2015-06-29 16:20:00	0.0	0.0	0.0	
1	2015-06-29 17:20:00	0.0	0.0	0.0	
2	2015-06-29 18:20:00	0.0	0.0	0.0	
3	2015-06-29 19:20:00	0.0	0.0	0.0	
4	2015-06-29 20:20:00	0.0	0.0	0.0	
...	...	...	...	...	
1223655	2022-06-30 19:20:00	0.0	0.0	0.0	
1223656	2022-06-30 20:20:00	0.0	0.0	0.0	
1223657	2022-06-30 21:20:00	0.0	0.0	0.0	
1223658	2022-06-30 22:20:00	0.0	0.0	0.0	
1223659	2022-06-30 23:20:00	0.0	0.0	0.0	

	precip2Day	precip3Day	precip7Day	precipMtd	precipYtd	\
0	NaN	NaN	NaN	NaN	NaN	
1	NaN	NaN	NaN	NaN	NaN	
2	NaN	NaN	NaN	NaN	NaN	
3	NaN	NaN	NaN	NaN	NaN	
4	NaN	NaN	NaN	NaN	NaN	
...	...	...	...	...	...	
1223655	0.0	0.0	2.0	9.0	238.0	
1223656	0.0	0.0	2.0	9.0	238.0	
1223657	0.0	0.0	2.0	9.0	238.0	
1223658	0.0	0.0	2.0	9.0	238.0	
1223659	0.0	0.0	2.0	9.0	238.0	

	pressureChange	...	uvIndex	visibility	windDirection	windGust	\
0	-1.4	...	2.0	16.09	NaN	NaN	
1	-1.0	...	1.0	16.09	NaN	NaN	

2	-0.3	...	0.0	16.09	NaN	NaN
3	0.3	...	0.0	16.09	NaN	NaN
4	0.9	...	0.0	16.09	NaN	NaN
...	...	...	...	...	...	...
1223655	2.7	...	0.0	13.55	110.0	NaN
1223656	3.7	...	0.0	13.59	80.0	NaN
1223657	3.8	...	0.0	13.85	80.0	NaN
1223658	3.1	...	0.0	13.17	50.0	NaN
1223659	2.0	...	0.0	9.51	50.0	NaN

	windSpeed	ID_ESTACION	año	mes	dia	hora
0	18.7	13	2015	6	29	16
1	18.0	13	2015	6	29	17
2	16.6	13	2015	6	29	18
3	15.1	13	2015	6	29	19
4	10.1	13	2015	6	29	20
...	...	...	...	...	...	...
1223655	10.8	8	2022	6	30	19
1223656	9.0	8	2022	6	30	20
1223657	7.9	8	2022	6	30	21
1223658	8.3	8	2022	6	30	22
1223659	10.4	8	2022	6	30	23

[1223660 rows x 37 columns]

```
[23]: encoded_id_estacion = meteo[meteo['mes'].isin([1,2,3,4,5,6])].
      ↪groupby(['ID_ESTACION'])[['uvIndex']].mean()
      encoded_id_estacion
```

```
[23]:          uvIndex
ID_ESTACION
0          1.361512
1          1.370795
2          1.391665
3          1.365001
4          1.380275
5          1.365692
6          1.353117
7          1.368457
8          1.366779
9          1.359866
10         1.362236
11         1.384423
12         1.361512
13         1.375568
14         1.366548
15         1.369807
```

```

16          1.371914
17          1.370828
18          1.364046
19          1.365956

```

```

[24]: #Mergeamos
df= df.merge(
    encoded_id_estacion,
    left_on='ID_ESTACION',
    right_index=True).sort_index()
df

```

```

[24]:
CAMPAÑA  ID_FINCA  ID_ZONA  ID_ESTACION  ALTITUD  VARIEDAD  MODO  \
4        14      14001      852           14  659.097938      81      1
5        14      17059      852           14  659.097938      81      1
6        14      87611      602           14  659.097938      81      1
7        14      12257      215           14  659.097938      17      1
8        14      97286      142           14  659.097938      17      1
...      ...      ...      ...      ...      ...      ...
9596     22      37461      239           6  700.000000      52      2
9597     22      58769      239           6  700.000000      32      2
9598     22      58769      239           6  700.000000      59      2
9599     22      88928      239           6  700.000000      40      2
9600     22      88928      239           6  700.000000      52      2

```

```

TIPO  COLOR  SUPERFICIE  PRODUCCION  windspeed_media  t_media  \
4      0      1      4.354864      14126.0      12.172118  12.094942
5      0      1      2.773533      6054.0      12.172118  12.094942
6      0      1      5.377166      12900.0      12.172118  12.094942
7      0      1      2.148042      5450.0      12.172118  12.094942
8      0      1      5.390867      30720.0      12.172118  12.094942
...      ...      ...      ...      ...      ...
9596     0      1      3.680000      NaN      11.664805  12.589099
9597     0      1      4.250000      NaN      11.664805  12.589099
9598     0      1      4.070000      NaN      11.664805  12.589099
9599     0      1      4.572700      NaN      11.664805  12.589099
9600     0      1      1.609900      NaN      11.664805  12.589099

```

```

rhum_media  precip_media  uvIndex
4      66.492972      40.389625  1.366548
5      66.492972      40.389625  1.366548
6      66.492972      40.389625  1.366548
7      66.492972      40.389625  1.366548
8      66.492972      40.389625  1.366548
...      ...      ...      ...
9596     66.438743      55.656785  1.353117
9597     66.438743      55.656785  1.353117

```

```

9598    66.438743    55.656785    1.353117
9599    66.438743    55.656785    1.353117
9600    66.438743    55.656785    1.353117

```

[9597 rows x 16 columns]

Ahora vamos a ver variable meteorológica está más correlacionada con la Produccion.

```

[25]: correlations = df.corr()['PRODUCCION'].drop('PRODUCCION')
sorted_correlations = correlations.abs().sort_values(ascending=False)
sorted_correlations

```

```

[25]: SUPERFICIE          0.758146
MODE                    0.270606
precip_media           0.181404
ALTITUD                0.146438
windspeed_media        0.131001
uvIndex                0.130138
t_media                0.117513
rhum_media              0.088291
ID_ESTACION            0.079442
ID_FINCA               0.042188
COLOR                  0.041029
VARIEDAD               0.035179
TIPO                   0.032299
ID_ZONA                0.032092
CAMPAÑA                0.027473
Name: PRODUCCION, dtype: float64

```

Vamos a usar F-test y mutual information para captar relaciones no lineales

```

[26]: from sklearn.feature_selection import f_regression, mutual_info_regression
import warnings
import matplotlib.pyplot as plt
warnings.filterwarnings("ignore")

```

```

[27]: # convertimos el DataFrame al formato necesario para scikit-learn
datos=df[df['CAMPAÑA'] != 22]
datos=datos.
      drop(['CAMPAÑA', 'ID_FINCA', 'ID_ZONA', 'ID_ESTACION', 'ALTITUD', 'VARIEDAD', 'MODE', 'TIPO', 'COLOR'])
datos

```

```

[27]: PRODUCCION  windspeed_media    t_media  rhum_media  precip_media  \
4          14126.0         12.172118  12.094942    66.492972    40.389625
5           6054.0         12.172118  12.094942    66.492972    40.389625
6          12900.0         12.172118  12.094942    66.492972    40.389625
7           5450.0         12.172118  12.094942    66.492972    40.389625
8          30720.0         12.172118  12.094942    66.492972    40.389625

```

...	...	...	...	...	...
8521	28160.1	11.664805	12.589099	66.438743	55.656785
8522	41310.0	11.664805	12.589099	66.438743	55.656785
8523	45420.0	11.664805	12.589099	66.438743	55.656785
8524	56140.0	11.664805	12.589099	66.438743	55.656785
8525	13869.9	11.664805	12.589099	66.438743	55.656785

	uvIndex
4	1.366548
5	1.366548
6	1.366548
7	1.366548
8	1.366548

...	...
8521	1.353117
8522	1.353117
8523	1.353117
8524	1.353117
8525	1.353117

[8522 rows x 6 columns]

```
[28]: data = datos.values
```

```
[29]: y = data[:,0:1]      # Produccion
      X = data[:,1:]      # el resto de variables meteorológicas

      feature_names = datos.columns[1:]

      # do calculations
      f_test, _ = f_regression(X, y)
      f_test /= np.max(f_test)


      mi = mutual_info_regression(X, y)
      mi /= np.max(mi)

      # do some plotting
      plt.figure(figsize=(20, 5))

      plt.subplot(1,2,1)
      plt.bar(range(X.shape[1]),f_test, align="center")
      plt.xticks(range(X.shape[1]),feature_names, rotation = 90)
      plt.xlabel('features')
      plt.ylabel('Ranking')
      plt.title('$F-Test$ score')
```

```
plt.subplot(1,2,2)
plt.bar(range(X.shape[1]),mi, align="center")
plt.xticks(range(X.shape[1]),feature_names, rotation = 90)
plt.xlabel('features')
plt.ylabel('Ranking')
plt.title('Mutual information score')

plt.show()
```



[30]: *#Vamos a utilizar para codificar ID\_ESTACION la variable precip\_media*

[31]: `df=df.  
↳drop(['t_media', 'rhum_media', 'uvIndex', 'ID_ESTACION', 'windspeed_media'],axis=1)`

[32]: `df`

```
[32]:
```

	CAMPAÑA	ID_FINCA	ID_ZONA	ALTITUD	VARIEDAD	MODO	TIPO	COLOR	\
4	14	14001	852	659.097938	81	1	0	1	
5	14	17059	852	659.097938	81	1	0	1	
6	14	87611	602	659.097938	81	1	0	1	
7	14	12257	215	659.097938	17	1	0	1	
8	14	97286	142	659.097938	17	1	0	1	
...	...	...	...	...	...	...	...	...	
9596	22	37461	239	700.000000	52	2	0	1	
9597	22	58769	239	700.000000	32	2	0	1	
9598	22	58769	239	700.000000	59	2	0	1	
9599	22	88928	239	700.000000	40	2	0	1	
9600	22	88928	239	700.000000	52	2	0	1	

	SUPERFICIE	PRODUCCION	precip_media
4	4.354864	14126.0	40.389625
5	2.773533	6054.0	40.389625
6	5.377166	12900.0	40.389625

7	2.148042	5450.0	40.389625
8	5.390867	30720.0	40.389625
...	...	...	...
9596	3.680000	NaN	55.656785
9597	4.250000	NaN	55.656785
9598	4.070000	NaN	55.656785
9599	4.572700	NaN	55.656785
9600	1.609900	NaN	55.656785

[9597 rows x 11 columns]

## 0.1 Estudio de Outliers

```
[33]: import matplotlib.pyplot as plt
import seaborn as sns
```

```
[34]: df14=df[df['CAMPAÑA']==14]
df15=df[df['CAMPAÑA']==15]
df16=df[df['CAMPAÑA']==16]
df17=df[df['CAMPAÑA']==17]
df18=df[df['CAMPAÑA']==18]
df19=df[df['CAMPAÑA']==19]
df20=df[df['CAMPAÑA']==20]
df21=df[df['CAMPAÑA']==21]
```

La metodología para el estudio de posibles outliers será la siguiente: - Dibujamos scatterplots para las variables que mayor correlación presentan con la variable a predecir (PRODUCCIÓN) en los distintos años. - Vemos si las observaciones que pueden ser consideradas outliers se mantienen todos los años o solo en años concretos. - Si se mantienen los mismos patrones todos los años, no las imputamos. - En caso contrario, las imputamos.

```
[35]: #Grafico x=SUPERFICIE
df14.plot(kind = 'scatter',x='SUPERFICIE',y = 'PRODUCCION')
plt.xlabel('# SUPERFICIE')
plt.ylabel('PRODUCCION')
plt.show()

df15.plot(kind = 'scatter',x='SUPERFICIE',y = 'PRODUCCION')
plt.xlabel('# SUPERFICIE')
plt.ylabel('PRODUCCION')
plt.show()

df16.plot(kind = 'scatter',x='SUPERFICIE',y = 'PRODUCCION')
plt.xlabel('# SUPERFICIE')
plt.ylabel('PRODUCCION')
plt.show()

df17.plot(kind = 'scatter',x='SUPERFICIE',y = 'PRODUCCION')
```

```

plt.xlabel('# SUPERFICIE')
plt.ylabel('PRODUCCION')
plt.show()


df18.plot(kind = 'scatter',x='SUPERFICIE',y = 'PRODUCCION')
plt.xlabel('# SUPERFICIE')
plt.ylabel('PRODUCCION')
plt.show()

df19.plot(kind = 'scatter',x='SUPERFICIE',y = 'PRODUCCION')
plt.xlabel('# SUPERFICIE')
plt.ylabel('PRODUCCION')
plt.show()


df20.plot(kind = 'scatter',x='SUPERFICIE',y = 'PRODUCCION')
plt.xlabel('# SUPERFICIE')
plt.ylabel('PRODUCCION')
plt.show()


df21.plot(kind = 'scatter',x='SUPERFICIE',y = 'PRODUCCION')
plt.xlabel('# SUPERFICIE')
plt.ylabel('PRODUCCION')
plt.show()


```























```
[36]: # Gráfico para x=precip_media
df14.plot(kind='scatter', x='precip_media', y='PRODUCCION')
plt.xlabel('precip_media')
plt.ylabel('PRODUCCION')
plt.show()

df15.plot(kind='scatter', x='precip_media', y='PRODUCCION')
plt.xlabel('precip_media')
plt.ylabel('PRODUCCION')
plt.show()

df16.plot(kind='scatter', x='precip_media', y='PRODUCCION')
plt.xlabel('precip_media')
plt.ylabel('PRODUCCION')
plt.show()

df17.plot(kind='scatter', x='precip_media', y='PRODUCCION')
plt.xlabel('precip_media')
plt.ylabel('PRODUCCION')
plt.show()
```

```


df18.plot(kind='scatter', x='precip_media', y='PRODUCCION')
plt.xlabel('precip_media')
plt.ylabel('PRODUCCION')
plt.show()

df19.plot(kind='scatter', x='precip_media', y='PRODUCCION')
plt.xlabel('precip_media')
plt.ylabel('PRODUCCION')
plt.show()


df20.plot(kind='scatter', x='precip_media', y='PRODUCCION')
plt.xlabel('precip_media')
plt.ylabel('PRODUCCION')
plt.show()


df21.plot(kind='scatter', x='precip_media', y='PRODUCCION')
plt.xlabel('precip_media')
plt.ylabel('PRODUCCION')
plt.show()


```























```
[37]: df14.plot(kind='scatter', x='MOD0', y='PRODUCCION')
plt.xlabel('MOD0')
plt.ylabel('PRODUCCION')
plt.show()

df15.plot(kind='scatter', x='MOD0', y='PRODUCCION')
plt.xlabel('MOD0')
plt.ylabel('PRODUCCION')
plt.show()

df16.plot(kind='scatter', x='MOD0', y='PRODUCCION')
plt.xlabel('MOD0')
plt.ylabel('PRODUCCION')
plt.show()

df17.plot(kind='scatter', x='MOD0', y='PRODUCCION')
plt.xlabel('MOD0')
plt.ylabel('PRODUCCION')
plt.show()

df18.plot(kind='scatter', x='MOD0', y='PRODUCCION')
```

```


plt.xlabel('MOD0')
plt.ylabel('PRODUCCION')
plt.show()

df19.plot(kind='scatter', x='MOD0', y='PRODUCCION')
plt.xlabel('MOD0')
plt.ylabel('PRODUCCION')
plt.show()


df20.plot(kind='scatter', x='MOD0', y='PRODUCCION')
plt.xlabel('MOD0')
plt.ylabel('PRODUCCION')
plt.show()


df21.plot(kind='scatter', x='MOD0', y='PRODUCCION')
plt.xlabel('MOD0')
plt.ylabel('PRODUCCION')
plt.show()


```























```
[38]: df14.plot(kind='scatter', x='ALTITUD', y='PRODUCCION')
plt.xlabel('ALTITUD')
plt.ylabel('PRODUCCION')
plt.show()

df15.plot(kind='scatter', x='ALTITUD', y='PRODUCCION')
plt.xlabel('ALTITUD')
plt.ylabel('PRODUCCION')
plt.show()

df16.plot(kind='scatter', x='ALTITUD', y='PRODUCCION')
plt.xlabel('ALTITUD')
plt.ylabel('PRODUCCION')
plt.show()

df17.plot(kind='scatter', x='ALTITUD', y='PRODUCCION')
plt.xlabel('ALTITUD')
plt.ylabel('PRODUCCION')
plt.show()


df18.plot(kind='scatter', x='ALTITUD', y='PRODUCCION')
```

```
plt.xlabel('ALTITUD')
plt.ylabel('PRODUCCION')
plt.show()


df19.plot(kind='scatter', x='ALTITUD', y='PRODUCCION')
plt.xlabel('ALTITUD')
plt.ylabel('PRODUCCION')
plt.show()


df20.plot(kind='scatter', x='ALTITUD', y='PRODUCCION')
plt.xlabel('ALTITUD')
plt.ylabel('PRODUCCION')
plt.show()


df21.plot(kind='scatter', x='ALTITUD', y='PRODUCCION')
plt.xlabel('ALTITUD')
plt.ylabel('PRODUCCION')
plt.show()
```























Como los patrones de las observaciones son similares para todos los años para las variables más correlacionadas con PRODUCCIÓN, decidimos no imputar ninguna observación

## 0.2 MODELOS

```
[39]: #Creamos un diccionario para ir guardando los resultados
results_test = {
    'XGboost': {'RMSE': None, 'MAE': None, 'R2': None},
    'GradientBoost': {'RMSE': None, 'MAE': None, 'R2': None},
    'Red Neuronal': {'RMSE': None, 'MAE': None, 'R2': None},
    'Bagging': {'RMSE': None, 'MAE': None, 'R2': None}
}
results_train = {
    'XGboost': {'RMSE': None, 'MAE': None, 'R2': None},
    'GradientBoots': {'RMSE': None, 'MAE': None, 'R2': None},
    'Bagging': {'RMSE': None, 'MAE': None, 'R2': None},
    'Red Neuronal': {'RMSE': None, 'MAE': None, 'R2': None}
}
```

```
[40]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error
```

```
from sklearn.metrics import r2_score
```

### 0.3 Conjunto de Entrenamiento

```
[41]: #Conjunto de datos de entrenamiento
df1421=df[df['CAMPAÑA'] != 22]
```

```
[42]: #Train y evaluación
X = df1421.drop('PRODUCCION', axis=1).values
y = df1421['PRODUCCION'].values.reshape(-1, 1)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

### 0.4 Gradient Boosting

```
[43]: # Creamos el modelo de Gradient Boosting
gbr = GradientBoostingRegressor(n_estimators=100, learning_rate=0.15,
↳max_depth=6, random_state=42)

# Ajustamos el modelo con los datos de entrenamiento
gbr.fit(X_train, y_train)

y_pred_test=gbr.predict(X_test)
```

```
[44]: # Vemos como ha sido el entrenamiento
y_pred_train=gbr.predict(X_train)
rmse_train = mean_squared_error(y_train, y_pred_train, squared=False)
mae_train = mean_absolute_error(y_train, y_pred_train)
r2 = r2_score(y_train, y_pred_train)
print(f'RMSE en train: {rmse_train:.5f}')
print(f'MAE en train: {mae_train:.5f}')
print("R2 score:", r2)

#Las guardo en el diccionario
results_train['GradientBoots']['RMSE'] = rmse_train
results_train['GradientBoots']['MAE'] = mae_train
results_train['GradientBoots']['R2'] = r2

# En test
rmse = mean_squared_error(y_test, y_pred_test, squared=False)
print('RMSE:', rmse)
mae=mean_absolute_error(y_test, y_pred_test)
print('MAE:', mae)
r2 = r2_score(y_test, y_pred_test)
print("R2 score:", r2)
```



```
#Las guardo en el diccionario
results_test['GradientBoost']['RMSE'] = rmse
results_test['GradientBoost']['MAE'] = mae
results_test['GradientBoost']['R2'] = r2
```

```
RMSE en train: 3261.51790
MAE en train: 2138.47911
R2 score: 0.9412560407833545
RMSE: 5282.798720503473
MAE: 3021.8006237884983
R2 score: 0.8202239812225536
```

## 0.5 XGBOOST

```
[45]: import xgboost as xgb
```

```
[46]: # Defino los hiperparámetros
params = {'objective': 'reg:squarederror',
          'max_depth': 6,
          'learning_rate': 0.1,
          'subsample': 0.5,
          'colsample_bytree': 0.5,
          'n_estimators': 100}

# Entreno el modelo
model = xgb.XGBRegressor(**params)
model.fit(X_train, y_train)

# Realizamos predicciones en el conjunto de prueba
y_pred_test = model.predict(X_test)
```

```
[47]: # Vemos como ha sido el entrenamiento
y_pred_train=model.predict(X_train)
rmse_train = mean_squared_error(y_train, y_pred_train, squared=False)
mae_train = mean_absolute_error(y_train, y_pred_train)
r2 = r2_score(y_train, y_pred_train)
print(f'RMSE en train: {rmse_train:.5f}')
print(f'MAE en train: {mae_train:.5f}')
print("R2 score:", r2)

#Las guardo en el diccionario
results_train['XGboost']['RMSE'] = rmse_train
results_train['XGboost']['MAE'] = mae_train
results_train['XGboost']['R2'] = r2

# En test
rmse = mean_squared_error(y_test, y_pred_test, squared=False)
```

```

print('RMSE:', rmse)
mae=mean_absolute_error(y_test, y_pred_test)
print('MAE:', mae)
r2 = r2_score(y_test, y_pred_test)
print("R2 score:", r2)

#Las guardo en el diccionario
results_test['XGboost']['RMSE'] = rmse
results_test['XGboost']['MAE'] = mae
results_test['XGboost']['R2'] = r2

```

RMSE en train: 4733.25977  
 MAE en train: 2848.65287  
 R2 score: 0.8762787323801671  
 RMSE: 5928.3168595662255  
 MAE: 3397.416068268864  
 R2 score: 0.7736052004736758

## 0.6 Bagging

```

[48]: from sklearn.ensemble import BaggingRegressor
      from sklearn.tree import DecisionTreeRegressor

```

```

[49]: # Creo modelo base de árbol de decisión
      base_model = DecisionTreeRegressor()

      # Creo modelo de Bagging con 100 estimadores
      bagging_model = BaggingRegressor(base_estimator=base_model, n_estimators=100,
      ↪random_state=42)

      # Entreno modelo con datos de entrenamiento
      bagging_model.fit(X_train, y_train)

      # Hago predicciones en conjunto de test
      y_pred_test = bagging_model.predict(X_test)

      # Vemos como ha sido el entrenamiento
      y_pred_train=bagging_model.predict(X_train)
      rmse_train = mean_squared_error(y_train, y_pred_train, squared=False)
      mae_train = mean_absolute_error(y_train, y_pred_train)
      r2 = r2_score(y_train, y_pred_train)
      print(f'RMSE en train: {rmse_train:.5f}')
      print(f'MAE en train: {mae_train:.5f}')
      print("R2 score:", r2)

      #Las guardo en el diccionario
      results_train['Bagging']['RMSE'] = rmse_train

```

```

results_train['Bagging']['MAE'] = mae_train
results_train['Bagging']['R2'] = r2

# En test
rmse = mean_squared_error(y_test, y_pred_test, squared=False)
print('RMSE:', rmse)
mae=mean_absolute_error(y_test, y_pred_test)
print('MAE:', mae)
r2 = r2_score(y_test, y_pred_test)
print("R2 score:", r2)

#Las guardo en el diccionario
results_test['Bagging']['RMSE'] = rmse
results_test['Bagging']['MAE'] = mae
results_test['Bagging']['R2'] = r2

```

RMSE en train: 2191.69715  
 MAE en train: 1124.53874  
 R2 score: 0.9734732058033906  
 RMSE: 5310.449504090665  
 MAE: 2874.3143787170675  
 R2 score: 0.8183371188555317

## 0.7 Red Neuronal

```

[50]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.metrics import RootMeanSquaredError
from tensorflow.keras.metrics import MeanAbsoluteError

# Definir la arquitectura de la red neuronal
model = Sequential()
model.add(Dense(64, activation='relu', input_shape=(X_train.shape[1],)))
model.add(Dense(64, activation='relu'))
model.add(Dense(1))

mae = MeanAbsoluteError()
rmse = RootMeanSquaredError()

# Compilar el modelo
model.compile(optimizer='adam', loss='mean_squared_error', metrics=[mae, rmse])

# Entrenar el modelo
model.fit(X_train, y_train, epochs=10, batch_size=32, validation_data=(X_test,
↪y_test))

```

```
# Evaluar el modelo en el conjunto de entrenamiento
loss, mae, rmse = model.evaluate(X_train, y_train)
r2 = 1 - (rmse**2 / tf.math.reduce_variance(y_train))
print("Pérdida en el conjunto de prueba:", loss)
print("RMSE en el conjunto de prueba", rmse)
print("MAE en el conjunto de prueba:", mae)
print("R2 en el conjunto de prueba:", r2)
```

Epoch 1/10

```
214/214 [=====] - 1s 2ms/step - loss: 204351648.0000 -
mean_absolute_error: 8463.8867 - root_mean_squared_error: 14295.1621 - val_loss:
166434384.0000 - val_mean_absolute_error: 7473.3296 -
val_root_mean_squared_error: 12900.9453
```

Epoch 2/10

```
214/214 [=====] - 0s 1ms/step - loss: 189596016.0000 -
mean_absolute_error: 8087.9746 - root_mean_squared_error: 13769.3867 - val_loss:
157651744.0000 - val_mean_absolute_error: 7893.4106 -
val_root_mean_squared_error: 12555.9443
```

Epoch 3/10

```
214/214 [=====] - 0s 1ms/step - loss: 184540800.0000 -
mean_absolute_error: 8196.5605 - root_mean_squared_error: 13584.5791 - val_loss:
155993008.0000 - val_mean_absolute_error: 7151.0796 -
val_root_mean_squared_error: 12489.7158
```

Epoch 4/10

```
214/214 [=====] - 0s 1ms/step - loss: 181883776.0000 -
mean_absolute_error: 8180.3018 - root_mean_squared_error: 13486.4297 - val_loss:
154250208.0000 - val_mean_absolute_error: 7666.9731 -
val_root_mean_squared_error: 12419.7510
```

Epoch 5/10

```
214/214 [=====] - 0s 1ms/step - loss: 181336768.0000 -
mean_absolute_error: 8203.1885 - root_mean_squared_error: 13466.1338 - val_loss:
156571600.0000 - val_mean_absolute_error: 8258.9648 -
val_root_mean_squared_error: 12512.8574
```

Epoch 6/10

```
214/214 [=====] - 0s 1ms/step - loss: 179982320.0000 -
mean_absolute_error: 8213.5850 - root_mean_squared_error: 13415.7490 - val_loss:
153430944.0000 - val_mean_absolute_error: 7567.6997 -
val_root_mean_squared_error: 12386.7246
```

Epoch 7/10

```
214/214 [=====] - 0s 1ms/step - loss: 179978672.0000 -
mean_absolute_error: 8248.8291 - root_mean_squared_error: 13415.6133 - val_loss:
153227456.0000 - val_mean_absolute_error: 7639.2065 -
val_root_mean_squared_error: 12378.5078
```

Epoch 8/10

```
214/214 [=====] - 0s 1ms/step - loss: 179719888.0000 -
mean_absolute_error: 8252.9473 - root_mean_squared_error: 13405.9648 - val_loss:
160659728.0000 - val_mean_absolute_error: 6931.7920 -
```

```

val_root_mean_squared_error: 12675.1621
Epoch 9/10
214/214 [=====] - 0s 1ms/step - loss: 180225152.0000 -
mean_absolute_error: 8188.8667 - root_mean_squared_error: 13424.7959 - val_loss:
152798464.0000 - val_mean_absolute_error: 7528.9785 -
val_root_mean_squared_error: 12361.1680
Epoch 10/10
214/214 [=====] - 0s 1ms/step - loss: 178962368.0000 -
mean_absolute_error: 8178.6772 - root_mean_squared_error: 13377.6816 - val_loss:
154488128.0000 - val_mean_absolute_error: 7181.6348 -
val_root_mean_squared_error: 12429.3252
214/214 [=====] - 0s 777us/step - loss: 181237472.0000
- mean_absolute_error: 7613.6636 - root_mean_squared_error: 13462.4463
Pérdida en el conjunto de prueba: 181237472.0
RMSE en el conjunto de prueba 13462.4462890625
MAE en el conjunto de prueba: 7613.66357421875
R2 en el conjunto de prueba: tf.Tensor(-0.0008561169972987059, shape=(),
dtype=float64)

```

```

[51]: #Las guardo en el diccionario
results_train['Red Neuronal']['RMSE'] = rmse
results_train['Red Neuronal']['MAE'] = mae
results_train['Red Neuronal']['R2'] = r2

```

```

[52]: # Evaluar el modelo en el conjunto de prueba
loss, mae, rmse = model.evaluate(X_test, y_test)
r2 = 1 - (rmse**2 / tf.math.reduce_variance(y_test))
print("Pérdida en el conjunto de prueba:", loss)
print("RMSE en el conjunto de prueba", rmse)
print("MAE en el conjunto de prueba:", mae)
print("R2 en el conjunto de prueba:", r2)

```

```

54/54 [=====] - 0s 866us/step - loss: 154488128.0000 -
mean_absolute_error: 7181.6348 - root_mean_squared_error: 12429.3252
Pérdida en el conjunto de prueba: 154488128.0
RMSE en el conjunto de prueba 12429.3251953125
MAE en el conjunto de prueba: 7181.634765625
R2 en el conjunto de prueba: tf.Tensor(0.004826661757593609, shape=(),
dtype=float64)

```

```

[53]: #Las guardo en el diccionario
results_test['Red Neuronal']['RMSE'] = rmse
results_test['Red Neuronal']['MAE'] = mae
results_test['Red Neuronal']['R2'] = r2

```

```

[54]: results_test

```

```
[54]: {'XGboost': {'RMSE': 5928.3168595662255,
  'MAE': 3397.416068268864,
  'R2': 0.7736052004736758},
  'GradientBoost': {'RMSE': 5282.798720503473,
  'MAE': 3021.8006237884983,
  'R2': 0.8202239812225536},
  'Red Neuronal': {'RMSE': 12429.3251953125,
  'MAE': 7181.634765625,
  'R2': <tf.Tensor: shape=(), dtype=float64, numpy=0.004826661757593609>},
  'Bagging': {'RMSE': 5310.449504090665,
  'MAE': 2874.3143787170675,
  'R2': 0.8183371188555317}}
```

```
[55]: results_train
```

```
[55]: {'XGboost': {'RMSE': 4733.259772344792,
  'MAE': 2848.6528656022,
  'R2': 0.8762787323801671},
  'GradientBoots': {'RMSE': 3261.517901237745,
  'MAE': 2138.47911297073,
  'R2': 0.9412560407833545},
  'Bagging': {'RMSE': 2191.697154089147,
  'MAE': 1124.5387358931494,
  'R2': 0.9734732058033906},
  'Red Neuronal': {'RMSE': 13462.4462890625,
  'MAE': 7613.66357421875,
  'R2': <tf.Tensor: shape=(), dtype=float64, numpy=-0.0008561169972987059>}}
```

## 0.8 PREDICCION

En base a los modelos que hemos entrenado, podemos ahora predecir la producción para la campaña 22. Consideramos que el modelo de Gradient Boosting es el más adecuado, por lo que lo emplearemos para dar una predicción.

```
[56]: # Seleccionamos las filas del dataset correspondientes a la campaña 22 y
      ↪ generamos con ellas
      # un dataframe
df22=df[df['CAMPAÑA']==22]
X = df22.drop('PRODUCCION', axis=1).values
```

```
[57]: # Generamos las predicciones
predicciones = gbr.predict(X)
```

```
[58]: # Las añadimos al dataframe
df22['PRODUCCION']=predicciones
```

```
[59]: # Vemos el resultado
df22
```

```
[59] :      CAMPAÑA  ID_FINCA  ID_ZONA  ALTITUD  VARIEDAD  MODO  TIPO  COLOR  \
8526      22      48626      302      600.0      32      2      0      1
8527      22      47921      302      600.0      32      2      0      1
8528      22      5696      919      655.0      59      1      0      1
8529      22      98814      919      655.0      32      2      0      1
8530      22      98814      919      655.0      40      2      0      1
...      ...      ...      ...      ...      ...      ...      ...
9596      22      37461      239      700.0      52      2      0      1
9597      22      58769      239      700.0      32      2      0      1
9598      22      58769      239      700.0      59      2      0      1
9599      22      88928      239      700.0      40      2      0      1
9600      22      88928      239      700.0      52      2      0      1
```

```
      SUPERFICIE  PRODUCCION  precip_media
8526      3.7503  21421.294747      39.370690
8527      7.3735  52235.319116      39.370690
8528      7.6200  4596.699651      40.389625
8529      3.3267  10427.578708      40.389625
8530      2.7724   9210.768702      40.389625
...      ...      ...      ...
9596      3.6800  27953.833890      55.656785
9597      4.2500  38946.035967      55.656785
9598      4.0700  40267.601216      55.656785
9599      4.5727  39914.968326      55.656785
9600      1.6099  11721.321435      55.656785
```

```
[1075 rows x 11 columns]
```